RESCIENCE C

Replication

[Re] Predicting Dynamic Embedding Trajectory in Temporal Interaction Networks

Romain Haton[1], Rayane Ait Ali Yahia[1], Vincent Gauthier[1] and Amel Bouzeghoub[1]
[1]SAMOVAR, Télécom SudParis, Institut Polytechnique de Paris, Palaiseau, France

Reproducibility Summary

Scope of Reproducibility – In this work, the article Predicting Dynamic Embedding Trajectory in Temporal Interaction Networks [1] is evaluated through a replication study. Replication results examine whether the claims made by authors are valid. Our goal is to replicate the experiments and achieve the same results as the authors.

Methodology – We reimplement the JODIE model of Kumar et al [1]. To test the authors' claims, we used the same architecture and made several modifications to the hyperparameters. We reproduce tables 3-4 for the authors' model of the original paper, excluding results from models in other papers. We extend the results of Figure 6 and the choice of embedding size from the original article

Result – After replication of the model and we validated most of the results found in the original paper, we have found only marginal difference between the replication and the original study. However, our extended experimentation about the robustness of embeddings size has lead us toward slightly different conclusion than those define in the original study. Additionally, we extended the original study by making some new experimentation about the impact of the `split` on the performance of the model, we have found that the `split` has a real impact on the overall performance.

What was easy – The authors had made their code available and all the necessary information at the following address https://github.com/srijankr/jodie. Their original code has considerably simplified our replication task.

What was difficult – Although the code is available. However some technical aspects have hindered our complete understanding of the model. We solved them by studying the code made available by the authors.

Communication with original authors – We contacted the authors for clarification about which aggregation method used to fit the inputs constraints of the RNN gate.

Introduction

Recommender systems are systems that cover a broad scope of techniques which all aim to provide suggestions of items that are most pertinent to a particular user (e.g., which music playlists to choose, which book a user might like, etc.), see [2] for a review of the techniques and challenges of the field. We can observe, particularly in areas such as e-commerce or social media, that one user may interact with several items/products, and these interactions may evolve with time and context (e.g., Winter clothes recommendations should be different from summer clothes recommendation). As a result, understanding and modeling the dynamic evolution of users and items is an important issue. In the literature, this dynamicity is generally modeled by an embedding trajectory in a Euclidean space. However, existing works suffer from several limitations, as the authors in [1] pointed out. At first, none of the existing methods predicts the future embedding trajectory. Instead, embeddings are generated only at each user's action. Moreover, static and dynamic properties are generally not considered in the same framework. Finally, scalability remains an open issue for the prediction of user-item interactions as well as for the model's training when dealing with large-scale datasets. To overcome these limitations, [1] proposes JODIE, a deep learning-based recommender system that models user-item interactions as a sequence of timestamped edges on a bipartite graph (cf. fig. 1). The core idea is that each user and item has two embeddings: a static embedding representing the entity's long-term stationary property and a dynamic embedding representing the time-varying property. These embeddings are used to make predictions regarding future interactions. The proposed method is trained with batches of data to overcome the scalability issue. The authors claim to outperform some of the state-of-the-art recommender systems such as Time-LSTM [3], LatentCross [4], Interaction Graph Embedding (IGE) [5], Recurrent Recommender Networks [6] and the Continuous-Time Dynamic Network Embeddings(CTDNE) [7] in predicting future interactions at least 20% better, and user state change predictions 12% better on average. This paper describes our efforts to replicate the JODIE model designed by Srijan Kumar et al [1]. We use the publicly available code provided by the authors to reproduce their results and validate their conclusions. Our code is available at https://github.com/ComplexNetTSP/JODIE.


Figure 1. Bipartite graph

Scope of Reproducibility

In this paper, we investigate the following claims from the original paper:

- **More accurate recommendation performance**. By using the JODIE model, the authors claim to achieve better predictions of future interactions over all the selected datasets. JODIE outperforms six models with an improvement of at least 20% and 14% in terms of Mean Reciprocal Rank (MRR) and Recall@10, respectively.

- **More accurate state changes performance**. The authors claim that JODIE outperforms five models regarding the user state change over all the selected datasets with an average improvement of 12.63% for the Area Under the Curve (AUC) metric.

- **Shorter learning time**. Thanks to the new t-batch algorithm, the learning phase is 9.2 times faster than other models comparable to JODIE, i.e., with a model composed of two Recurrent Neural Networks (RNNs).

- **Robustness of the model to the proportion of the training set**. Regardless of the training data percentage, JODIE outperforms all the compared models.

- **Robustness of the model to the embedding size**. The embedding size does not have much impact on the performance of JODIE.

We have implemented the JODIE model using a more recent version of the Pytorch library aiming at reproducing the results stated by the authors:

- We explored which embedding size is the most effective for predicting state change and future interaction.

- We also replicated the experiments that concern the robustness of the JODIE model to the percentage of training data.

- To measure the robustness to the dynamic embedding size, the authors measured the performance of JODIE with embedding sizes from 32 to 256 on solely the LastFM dataset. We extended this experiment by adding three embedding sizes, 8, 16, and 32, on the LastFM and Wikipedia datasets and by calculating the MRR (Mean Reciprocal Rank) and Recall@10.

- In addition to reproducing the results presented in the paper, we perform a novel experiment that tests the impact of the hyperparameter split. We varied split by 5, 500, and 50000 on the MOOC dataset, which will increase the number of backpropagation, and we calculate the performance in AUC.

The main issue we encountered during the replication of the original model [1] was the difficulty in understanding the t-batch algorithm, as the original article did not provide nor reference it. Still, we managed to find that the algorithm had already been published by the same authors as a *preprint* in [8]. To sum up, we have reproduced the results of the original paper by extending the experiment on the robustness of the dynamic embedding size and by bringing an additional experiment regarding the choice of the hyperparameters.

Methodology

JODIE is a model that learns dynamic embeddings of users $u_t \in \mathbb{R}^n$, $\forall u \in \mathcal{U}$ and items $i_t \in \mathbb{R}^n$, $\forall i \in \mathcal{I}$ over time t, $\forall t \in [0; T]$, where $\mathcal{U}$ and $\mathcal{I}$ are the sets of users and items and T the final time. Moreover, the interactions are ordered in time, i.e., an interaction between a user u_r and an item i_r at time t_r characterized by f_r feature vector is noted $S_r = (u_r, i_r, t_r, f_r)$. Figure 2 shows the structure of JODIE. The model comprises two main components: the update operator, which, with the help of two RNNs, updates the embeddings, and the projection operator, which forecasts the embedding into a future time, represented in purple in the figure 2. Then, the model uses the different outputs to predict the future interaction or the state change.

Figure 2. Pipeline. MLP_{reg} gives an estimation of the embedding of the next item $\widetilde{i}$ with which a user u is likely to interact. The MLP is designed to recommend the best i to each user. $MLP_{classif}$ is used to predict the user's state change $\widehat{u}_{state}$. Will the user potentially drop out / get banned ?

Code

We re-implemented the JODIE model using newer version of Pytorch [9] version 1.10 and python 3.8 independently of the original code provided by the authors[1]. We want to highlight the fact that we used the *tanh* function as the activation function at the output of the RNNs (see fig. 2) to build our model, in accordance to the authors' original code but contrary to what's written in the original paper. Our code is available at the following address: https://github.com/ComplexNetTSP/JODIE.
Our source code is organized as follows:

- **preprocessing.py** where everything about the data preprocessing and the t-batch algorithm.

- **model.py** the model as described in figure 2.

- **train.py** is used to train the model.

- **evaluate.py** is used to evaluate the model after training.

By organizing the work in this way, it allows to have a modular, easily reusable and understandable code. To run the model, it's necessary to create an environment with the required packages and run one of the following python scripts in order to lauch the trainning/evaluation of the JODIE model on one of avaialble dataset:

- *train_evaluate_wikipedia.py, train_evaluate_lastfm.py, train_evaluate_reddit.py*

- *train_evaluate_reddit_state.py, train_evaluate_wikipedia_state.py, train_evaluate_mooc_state.py*

More informations are available in README file in the repositoritory of our code[2].

Model descriptions

Update operator – The update operator takes the form of two mutually recursive neural networks shown in Figure 3. The RNN$_U$ is used to update the dynamic user embeddings and the RNN$_I$ updates the dynamic item embeddings. In this case, the embeddings are considered as the hidden state of the classical RNNs respectively. At time $t = 0$, the embeddings are initialized randomly according to uniform distribution $\mathcal{U}[0; 1[$ followed

---

[1]https://github.com/srijankr/jodie
[2]https://github.com/ComplexNetTSP/JODIE

by a normalization and are noted u_0 and i_0. As input, there is the embeddings at time $t = 0$ but also Δ the time elapsed between an entity and the previous interaction and f a feature vector which characterizes the link of the interaction. A specificity of these RNNs is that they take four inputs instead of two as usual. In order to narrow down to just two inputs, we concatenate the embedding of the other entity, Δ and f. In Figure 3, we denote the concatenation by [.,.]. More formally, we use the following formulas to update the embeddings:

$$u^+ = \sigma \left(W_1^u u^- + W_2^u i^- + W_3^u f + W_4^u \Delta_u \right)$$

$$i^+ = \sigma \left(W_1^i i^- + W_2^i u^- + W_3^i f + W_4^i \Delta_i \right)$$

Where the matrices $W_1^e, ..., W_4^e$ are the parameters of the RNN_e and σ an activation function (here hyperbolic tangent). We note by u^- and i^- the embeddings before update and by u^+ and i^+ the embeddings after update. Once this step is finished, we can proceed to the second operation, the projection operation.
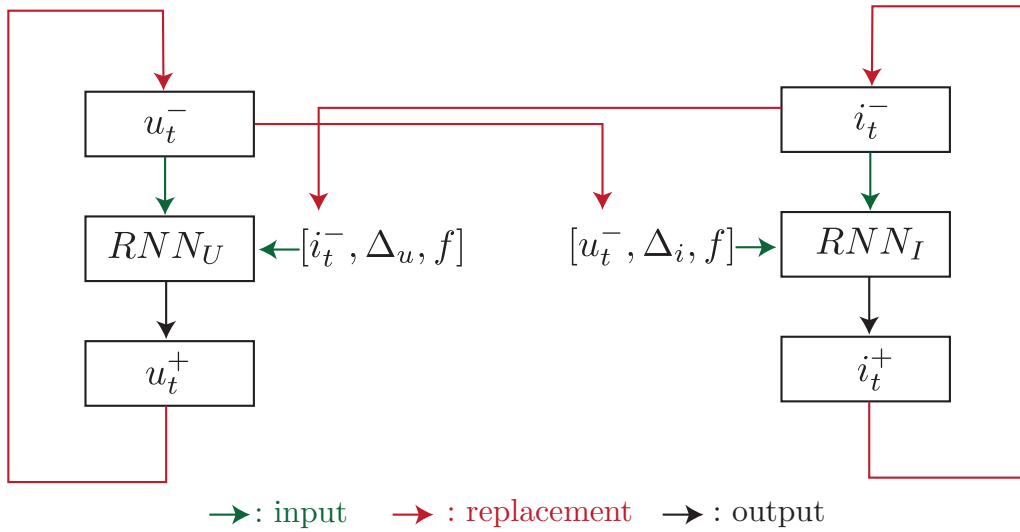

$\longrightarrow$: input $\longrightarrow$: replacement $\longrightarrow$: output

Figure 3. RNN mutually recursive

Projection operation – This operation consists in projecting the user's embedding to a future time. Indeed, it is assumed that user's interest changes with time, for example in winter and in summer, hence the need to model this variation as a linear projection. To a more deeper understanding please refer to the orginal figure provide by the authors themself in figure 3 in [1]).

We notice that over a very short time interval, the projected embedding of the user changes slightly, but as time goes, the projected embedding will be further away from its original embedding. To perform this operation, we must first convert Δ into a vector $w \in \mathbb{R}^n$ using a linear layer represented by its weights vector W_p. We have $w = W_p \Delta$. We initialize W_p by a Gaussian of zero mean. We obtain the projected embedding by computing Hadamard product, denoted $*$, of $1 + w$ and u_t. The authors use the following formula.

$$\widehat{u}_{t+\Delta} = (1 + w) * u_t$$

We can see that if $\Delta = 0$ then $w = 0$ and the projected embedding is the same as the initial embedding. We can consider this operation as a simple linear transformation.

These two operators consist in the learning of the model. One limitation cited by the authors is the learning time. Indeed, the existing models process each interaction one after the other, which takes a considerable amount of time. The authors have invented

an algorithm that allows to separate temporal data in batches and thus to accelerate the learning time. This algorithm is called `t-batch`.

t-Batch algorithm – The `t-batch` algorithm is a linear algorithm that split temporal data into batches (see fig. 4). This algorithm is a key point of the method that allows to speed up the learning time considerably. The `t-batch` algorithm is explained in the original article only by a few lines. Regarding the code, the algorithm was difficult to identify since it was mixed with other parts of code in the learning loop. To facilitate the understanding of the `t-batch` we propose the algorithm 1. To summarize the 1 algorithm, given a list of timestamped events, it dispatches the event into a batch according to the following constrains first it ensures that the same batch cannot contain the same entity multiple times. Secondly, it ensures to retain temporal dependencies inside batches and between batches, i.e., we process interactions in order. By meeting these two conditions, we guarantee that batches are independent (and parallelizable). Finally, in line 9. of Alg. 1 we compute the index of the batch B_{idx} where the interaction $\mathcal{S}_r \in \mathcal{S}$ will be put in. Let maxBatch$(e, r-1)$ be the function that gives the index of the batch at time $r-1$ containing the entity e.

Algorithm 1 t-Batch

1: **Input** : Sequence of interactions ordered by time $S : S_j = (u_j, i_j, t_j, f_j)$
2: **Output** : Sequence of batches $B : B_k = \{S_{k,1}, ..., S_{k,n}\}$
3: **Initialize** :
4:     tbatch_id_u[u] $\leftarrow 0 \ \forall \ u \in \mathcal{U}$,   tbatch_id_i[i] $\leftarrow 0 \ \forall \ i \in \mathcal{I}$
5:     $B_k \leftarrow \{\}, \ \forall \ k \in [\![1, \ \text{Card}(S)]\!]$
6:     $C \leftarrow 0$
7: **for** $S_j \in S$ **do**
8:     Extraction $S_j = (u_j, i_j, t_j, f_j)$
9:     idx $\leftarrow$ max(tbatch_id_u[u_j], tbatch_id_i[i_j]) $+ 1$
10:     $B_{\text{idx}} \leftarrow B_{\text{idx}} \cup \{S_j\}$
11:     tbatch_id_u[u_j] $=$ idx
12:     tbatch_id_i[i_j] $=$ idx
13:     $C \leftarrow \max(C, \text{idx})$
14: **end for**
15: **return** $\{B_1, ..., B_C\}$

Next item embedding and state change losses – The JODIE model has been designed to give directly an embedding of an item, which is time saving. Indeed, the existing models predict a probability for each item. This can be very long if the dataset contains many items. JODIE will produce an embedding and the suggested item will be the one whose embedding will be the closest to the one of the prediction. To make this prediction, we use a neural network that will have as input the projected embedding of the user at time $t + \Delta$ noted $\widehat{u}(t + \Delta)$ and the embedding of the previous item of the user before time $t+\Delta$ noted $i(t+\Delta^-)$. This embedding is important because item can interact with other users between time t and $t + \Delta$, which means that it contains more recent information. JODIE uses both static and dynamic embeddings, denoted respectively $\bar{e}$ and e, to make the prediction of a static and dynamic embedding $\widetilde{i}$. The prediction is done with a fully connected linear layer.

$$\widetilde{i}(t + \Delta) = W_1\widehat{u}(t + \Delta) + W_2\bar{u} + W_3 i(t + \Delta^-) + W_4\bar{i} + B$$

Where $W_1, ..., W_4$ and the bias vector B are the model weights.

JODIE is trained to minimize the L2 distance between the predicted embedding and the real item embedding of each interaction. We compute the loss function as follows:

$$Loss = \sum_{(u,\,i,\,t,\,f)\in S} \| \widetilde{i}_t - [\bar{i}, \, i_t^-]\|_2 + \lambda_U \|u_t - u_t^-\|_2 + \lambda_I \|i_t - i_t^-\|_2 \tag{1}$$

The first term minimizes the error of the predicted embedding. The next twos terms regulate the loss function and allow the dynamic embeddings of users and items to not vary too much. λ_U and λ_I are the regularization terms that penalize the two terms. In the original paper, the authors fixed the values of λ_U and λ_I as: $\lambda_U = \lambda_I = 1$, without much justification. It is not clear if this is the best value or if another value would be more appropriate.

When we want to predict the change of state of a user, we use the same loss function with an additional term which is the cross-entropy. To do this, we have to make sure that the classes are binary. In this case, it is possible to train the model using an additional loss term function that will make it able to predict the labels using the user's embedding after an interaction. The additional loss function term, for predicting a user's state change, is a weighted cross-entropy expressed in the following equation:

$$\ell(x, y) = \sum_{n=1}^{N} \frac{l_n}{\sum_{n=1}^{N} w_{y_n}} \ \text{ with } \ l_n = -w_{y_n} \log\left(\frac{\exp(x_{n,y_n})}{\sum_{c=1}^{C} \exp(x_{n,c})}\right) \tag{2}$$

rite the corresponding paper using the proposed LaTeX template or your own t Where x is the input, y is the target, w is the weight, C is the number of classes (here two) and N is the batch size. The loss function becomes:

$$Loss = \sum_{(u,\,i,\,t,\,f)\in S} \| \widetilde{i}_t - [\bar{i}, \, i_t^-]\|_2 + \lambda_U \|u_t - u_t^-\|_2 + \lambda_I \|i_t - i_t^-\|_2 + \ell(u_{true}, \widehat{u}_{pred}) \tag{3}$$

Where u_{true} is the real class of the user and $\widehat{u}_{pred}$ his predicted class. It is also worth noting that the regularizations terms are MSE. The authors wanted the embeddings not to vary heavily between the previous time $t-1$ and the current time t. Indeed, they started from the postulate that the behavior of a user represented by its dynamic embedding does not vary considerably over a small temporal interval.

Training JODIE model – We detail the learning steps and the specificity of JODIE model we provide.

1. We use a function called **preprocess** which is in the file **preprocessing.py** to make a pre-processing of the data to extract information necessary during the training. Information like the sequence of users, items, time, features or previous items that will be used during the learning step.

2. We initialize each embedding randomly $\mathcal{U}([0, 1[)$.

3. During the trainning phase we usually we split the observations from the original dataset within a given number of batches (K times periods, see fig. 4 for more detail). But when we have a dataset has a temporal dependency, we cannot operate under the same assumption, that why in the JODIE model they created the `t-batch` algorithm. First, we split evenly the dataset into ordered batch of data secondly, we apply the `t-batch` algorithm independently on each subset of the data as described in the figure 4. Our implementation of the `t-batch` differ from what was originally defined in [8], in our implementation do all the t-batches at once with the function **t_batch** in the file **preprocessing.py** before the learning step and loop over the periods.

4. During the learning stage, in the function **train_ray** in the file **train.py**, we use the function **projection** from the file **model.py** to project users' embedding in a future time. We also use the function **predict_embedding_item** in the file **model.py** to make a prediction of the next embedding of the item. And finally, we calculate the Mean Square Error (MSE) of the predicted embedding and the real embedding.

5. We update embeddings of users and items using the functions **update_rnn_user** and **update_rnn_item** respectively in the file **model.py**.

6. We compute the two regularization terms using the function **regularizer** of the file **preprocessing.py**. This function uses the hyperparameter λ_e.

7. If necessary, we predict the state prediction and we calculate the loss Cross Entropy (CE) with the function **loss_predict_state** of the file **model.py** which makes the prediction and the calculation of the loss.

8. We perform the back propagation of the gradient, at the end of a period, to update the model weights.

9. Once all the time periods are covered, we go to the second epoch.

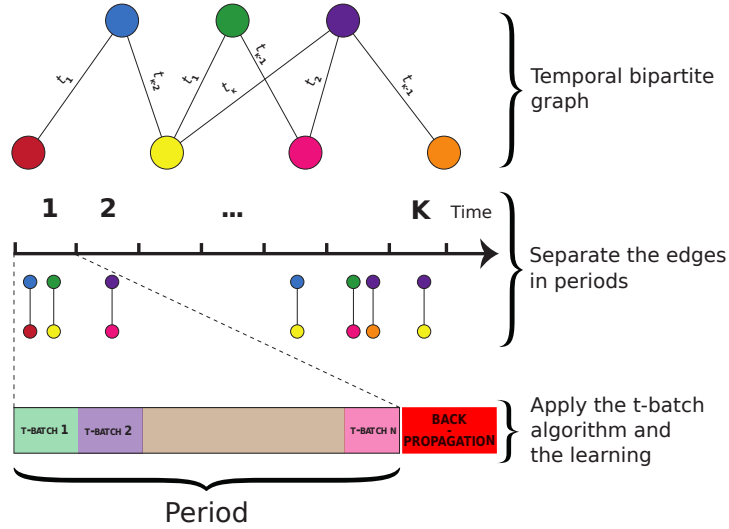Once the learning stage is over, the evaluation is carried out.

Figure 4. Pipeline of the `t-batch` algorithm

Evaluate JODIE model – The prediction and evaluation stage are just as special as the learning stage. In the evaluation, there is still a "learning" because during the predictions, the model will update itself respecting the time of the periods defined by `split`. This evaluation is done thanks to the function **evaluate** of the file **evaluate.py**. During the evaluation, we do not apply the `t-batch` algorithm on the remaining interactions, but we will treat them one after the other (figure 4 not t-batch but interaction just before the back-propagation).
From a temporal bipartite graph, we proceed in the same way as for learning, we separate the edges in temporal order. Then, we will separate the data into periods and we will process the interactions one after the other. The periods are defined with the hyper-parameter `split` as for the periods of the `t-batch`. In addition to all the learning steps in the evaluation, we add a prediction step. According to the prediction task, we will predict a probability of a user's state change using the function **predict_state** in the file **model.py**. Otherwise we will predict an embedding and look for the item which

has its closest embedding in terms of Euclidean distance. Finally, we will calculate the metrics to obtain the validation and generalization error.

Datasets

The data sets that were used in the original article were also used in this work. The authors of the article give the datasets.

Reddit [10] post dataset: this public dataset consists of one month of posts made by users on subreddits. We selected the 1000 most active subreddits as items and the 10000 most active users. We convert the text of each post into a feature vector representing their Linguistic Inquiry and Work Count [11] (LIWC) categories.
Wikipedia [12] edits: this public dataset is one month of edits made by edits on Wikipedia pages. We selected the 1000 most edited pages as items and editors who made at least 5 edits as users (a total of 8227 users). Similar to the Reddit dataset, we convert the edit text into a LIWC feature vector.
LastFM [13] song listens: this public dataset has one month of who-listens-to-which song information. We selected all 1000 users and the 1000 most listened songs. In the dataset, interactions do not have features.
Reddit bans: Reddit post dataset with ground-truth labels of banned users from Reddit. This gives 366 true labels (=0.05%).
Wikipedia bans: Wikipedia edit data with public ground-truth labels of banned users. This results in 217 positive labels (=0.14%).
MOOC [14] student drop-out: this public dataset consists of actions done by students on a MOOC online course. This dataset consists of 7047 users interacting with 98 items. There are 4066 drop-out events (=0.98%). All datasets and their details are shown in table 1.

	Users	Items	Interactions	State changes	Action repetition	Features size
Reddit	10,000	984	672,447	366	79%	172
Wikipedia	8,227	1,000	157,474	217	61%	172
LastFM	980	1,000	1,293,103	-	8.6%	2
MOOC	7,047	97	411,749	4,066	-	4

Table 1. Description of the datasets that were used in this project

The authors of the original paper created their own datasets. Thus, a dataset has the following format :

- A line represents an interaction or an edge

- Each line is described by a user, an item, a timestamp, a state label and features

- The timestamp is in cardinal number format

- The state label is 1 if the user changes state and 0 otherwise. If there are no state labels, use 0 for all interactions

- Features can be as long as desired and of dimension at least 1. If there is no feature, use 0 for all interaction

A negative point about the data sets is that the variable names are not in the right places and this creates errors when using them. That is why we created a function called **fetch_datasets** in the file **preprocessing.py** which allows to have correct variable names

which have the following format: `user, item, timestamp, labels, f_1, ..., f_n` with n the number of features.

Hyperparameters

As described in the model descriptions section, the model has several hyperparameters like epoch number, dynamic embedding size, learning rate, `split`, λ_U and λ_I. Initially the hyperparameters were chosen as in the original article as follow.

- Embedding size: 32, 64, **128**, 256

- Epoch, learning rate, λ_U and λ_I: 50, $1e - 3$, 1, 1 respectively

- `Split`: 500

Only the size of the dynamic embeddings have been studied and chosen using a grid search. The other hyperparameters are just given without any justification.

Extended experiments

Of all these hyper-parameters, the one that caught our eye the most was `split`. Indeed, `split` is used to define the frequency at which t-batches are created and the JODIE model is trained. It is important to choose its value well to have enough data in a period. To test this hyper-parameter, we set the other hyper-parameters to 50 epoch, $\lambda_U = \lambda_I = 1$, learning rate $= 1e - 3$ and the size of the embeddings to 128, and we test these 5, 500 and 50000 for the value of `split`. We define `split` as follows: `split` is used to split the dataset for a specific period of time. The authors use it as follows : $\frac{time_{end} - time_{start}}{split}$

Results

In this section, we will present the results obtained with replication. First, we tried to replicate the results for the user's state change prediction on all available datasets. The strategy was to predict only with the model weights obtained in the last learning epoch while the authors decided to keep the best epochs weights. For this prediction task, other embedding sizes were tested like 8, 16 and 32. Then, we tried to reproduce the results for the future interaction prediction using the same strategy as for the state prediction. Then a second part was devoted to adding results that are not present in the original paper like the importance of the `split` hyper-parameter.

Core replication results

The first part of the results concerns the prediction of the user's state change. The model has for hyperparameters a `split` of 500, $\lambda_U = \lambda_I = 1$ as those in the original paper and the embedding size varies. We present the results obtained for embeddings of size 32, 128 and 256 as well as the results of the original paper in the table 2.
We observe for the Reddit dataset, a performance of 0.586 AUC for the reproduction against 0.599 in the original paper. We also notice a 0.718 AUC for the MOOC dataset results which is close to the 0.756 AUC obtained in the original paper. For Wikipedia dataset, we can observe a performance of 0.738 AUC for the reproduction and 0.831 AUC for the original model. On the Reddit, Wikipedia and MOOC datasets, we see a standard deviation of 0.018, 0.030 and 0.021 respectively for an embedding size of 128. We can explain this difference by the fact that the authors of JODIE evaluated their models on each epoch and chose the best performance in validation while our results come from the 50 and last epoch. This can explain the difference between their results and our results

	JODIE	Replications		
	128	32	128	256
Reddit	**0.599**	0.563 ± 0.013	0.586 ± 0.018	0.577 ± 0.012
Wikipedia	**0.831**	0.728 ± 0.010	0.738 ± 0.030	0.711 ± 0.008
MOOC	**0.756**	0.688 ± 0.006	0.718 ± 0.021	0.711 ± 0.004

Table 2. Results: user state change prediction

but this difference remains nevertheless marginal. We can see that if we increase the size of the embeddings to 256, the results degrade slightly. This can be due to size of the embeddings which is too big which allows fitting some noise from the data. Conversely, if the size of the embeddings is 32, the results are worse. This can be explained by the fact that the embeddings are too small and lack information to make a good prediction.

For the prediction of future interaction, we used the same hyperparameters as for the state prediction. We present the obtained results and the results of the original paper in the table 3

Dataset	JODIE	Replications		
	128	32	128	256
Reddit				
MRR	**0.726**	0.711 ± 0.007	0.719 ± 0.008	0.684 ± 0.005
Recall@10	**0.852**	0.810 ± 0.020	0.830 ± 0.020	0.747 ± 0.011
Wikipedia				
MRR	0.746	$\mathbf{0.763} \pm 0.005$	$\mathbf{0.764} \pm 0.002$	$\mathbf{0.763} \pm 0.004$
Recall@10	0.822	$\mathbf{0.829} \pm 0.005$	$\mathbf{0.833} \pm 0.004$	$\mathbf{0.827} \pm 0.004$
LastFM				
MRR	0.195	$\mathbf{0.311} \pm 0.001$	$\mathbf{0.313} \pm 0.002$	$\mathbf{0.312} \pm 0.001$
Recall@10	0.307	$\mathbf{0.455} \pm 0.003$	$\mathbf{0.483} \pm 0.047$	$\mathbf{0.453} \pm 0.002$

Table 3. Results: future interaction prediction

We notice that only the results on the Reddit dataset are slightly lower than in the original paper. On these data, we obtained 0.719 in MRR and 0.830 in Recall@10. However, these results remain close to those obtained by the authors with 0.726 in MRR and 0.852 in Recall@10. For the Wikipedia and LastFM datasets, we can see that our results outperformed the authors' results with 0.764 in MRR and 0.833 in Recall@10 and 0.313 in MRR and 0.483 in Recall@10 respectively. These differences can be explained by the choice to evaluate the last epoch and not to choose the best performance on the validation, nonetheless these differences remain small.

To go further in the replication and reproduction, we will also reproduce two figures that show that the robustness of JODIE model (see fig 5). The figure 5 A) will be on the same dataset but for the prediction of user state change. For this, we varied the percentage of the training set from 20% to 60% in 10% steps. For the state changes prediction in figure 5 A) shows that due to class imbalances, the larger the training is, less the model is accurate in predicting state change.

Figure 5. Robustness of JODIE replication: the figure B) compare the MRR of JODIE replication with baselines on interaction prediction task, by varying the training data size. The figure A) shows the AUC of user state change prediction task by varying the training data size.

The figure 5 B) will concern the prediction of future interaction on the Wikipedia dataset by plotting the MRR. For this figure, we varied the percentage of the training set from 10% to 80% with 10% step. On the contrary with the interaction prediction as shown in the figure 5 B), larger the training set is, the better is the model will be able to predicting the future interactions. This describes the classical behavior, where the more training data will lead to better results (similar observation was made in the original paper).

Additionally, we tested the impact of the embedding size on the model. In figure 6 A) shows the MRR results of predicting future interactions for the datasets LastFM and Wikipedia and in the figure 6 B), we show the same results but with the recall@10. Despite the authors claim that the overall performance doesn't depend on the embedding size, in figure 6, we show that depending on the metric of interest (i.e.: MRR or Recall@10) the embedding size has an impact on the overall performance of the model.

As conclusion, after having reproduced their results we also claim that their model outperforms the existing models on these 4 datasets and that the replication and reproduction are validated.
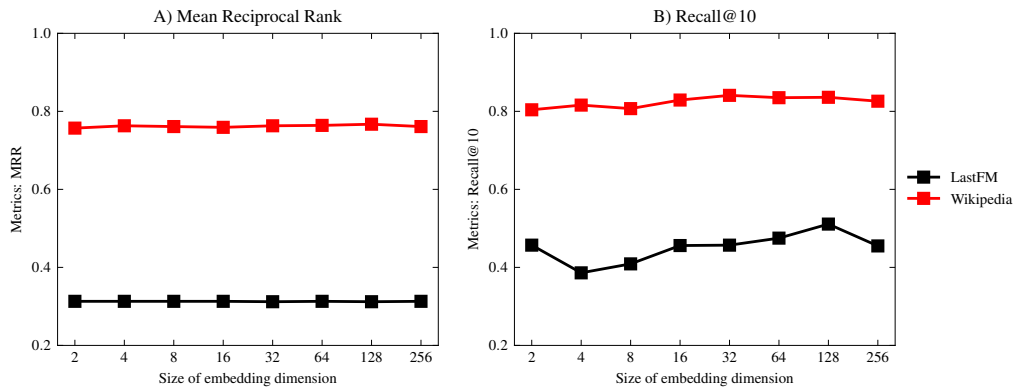

Figure 6. Robustness of JODIE replication

Additional result

We decided extend our analysis of the original model by testing the impact of the split (see sec. Training JODIE model) on the performance of the JODIE model, as a reminder,

the original study doesn't consider `split` as a hyper-parameter but as was a constant fixed at 500. In the figures 7, we show that this parameter has a strong importance on the overall results. We can see in figure 7, for any embedding size, that the performance is always increasing when the `split` value is higher. By increasing this `split` value from 5 to 500, we have increased the performance by 0.14 in AUC for the different sizes of embedding. Going from 500 to 50000, increased the performance by 0.08. We can explain this improvement by the fact that the model will do back-propagation after each time step defined by `split`. So by increasing the `split` number, we will increase the number of back-propagation. This suggests that higher values of this parameter allow to enhance the performance but the choice of it must be done considering the running time and the complexity which are increasing significantly when the `split` get larger.
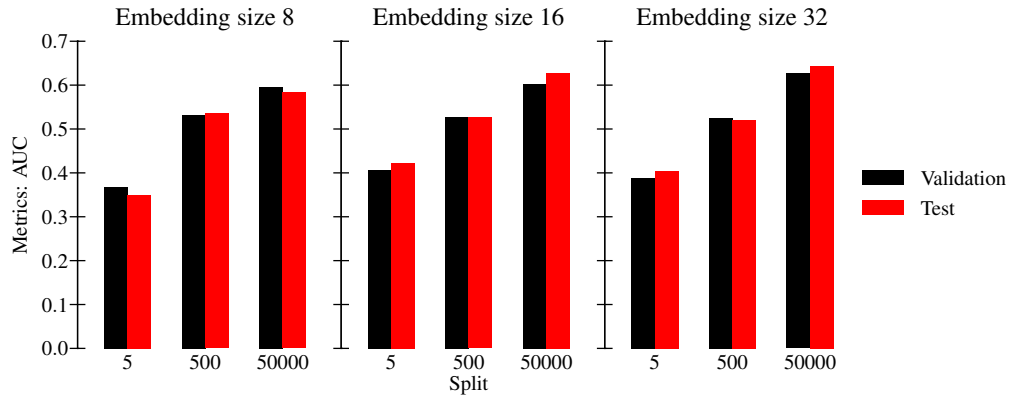
Figure 7. Comparison the robustness of the JODIE model as a function of the hyper-parameter `split`.

Conclusion

We are able to successfully replicate the original code provided by the authors and reproduce the results obtained in original JODIE publication. Even if our reproduction more or less match the result provided in the original study, we have found some sligh difference compare with the result in the original paper. This can be explained by a variation due to the initialization. However, for LastFM case our reproduction result clearly outperform those claimed in the original study, we are suspecting some reporting error in the original study. Secondly, we have show that the impact of the embedding size has a slight impact on the result if a different evaluation metric than the one proposed in the original paper was used. This fact balance the claim made by the authors stating the embedding size has no impact on the overall performance of their model. Finally, we have found that the split has a strong impact on the overall performance, This hyper-parameter unexplored in the original paper must be adjusted taking into account a complexity-performance trade-off of the model. As perspective, we think it would be interesting to use the focal loss [15] instead of the standard cross entropy loss in the JODIE model in order to alleviate the issue we have met with unbalanced class.

References

1. S. Kumar, X. Zhang, and J. Leskovec. "Predicting Dynamic Embedding Trajectory in Temporal Interaction Networks." In: **Proceedings of the 25th ACM SIGKDD international conference on Knowledge discovery and data mining**. ACM. 2019.

2. F. Ricci, L. Rokach, and B. Shapira. **Recommender Systems: Techniques, Applications, and Challenges**. Springer US, Oct. 2021, pp. 1–35.

3. Y. Zhu, H. Li, Y. Liao, B. Wang, Z. Guan, H. Liu, and D. Cai. "What to Do Next: Modeling User Behaviors by Time-LSTM." In: **Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence, IJCAI-17**. 2017, pp. 3602–3608.

4. A. Beutel, P. Covington, S. Jain, C. Xu, J. Li, V. Gatto, and E. H. Chi. "Latent Cross: Making Use of Context in Recurrent Recommender Systems." In: **Proceedings of the Eleventh ACM International Conference on Web Search and Data Mining**. WSDM '18. Marina Del Rey, CA, USA: Association for Computing Machinery, 2018, pp. 46–54.

5. Y. Zhang, Y. Xiong, X. Kong, and Y. Zhu. "Learning Node Embeddings in Interaction Graphs." In: **Proceedings of the 2017 ACM on Conference on Information and Knowledge Management**. CIKM '17. Singapore, Singapore: Association for Computing Machinery, 2017, pp. 397–406.

6. C.-Y. Wu, A. Ahmed, A. Beutel, A. J. Smola, and H. Jing. "Recurrent Recommender Networks." In: **Proceedings of the Tenth ACM International Conference on Web Search and Data Mining**. WSDM '17. Cambridge, United Kingdom: Association for Computing Machinery, 2017, pp. 495–503.

7. G. H. Nguyen, J. B. Lee, R. A. Rossi, N. K. Ahmed, E. Koh, and S. Kim. "Continuous-Time Dynamic Network Embeddings." In: **Companion Proceedings of the The Web Conference 2018**. WWW '18. Lyon, France: International World Wide Web Conferences Steering Committee, 2018, pp. 969–976.

8. S. Kumar, X. Zhang, and J. Leskovec. "Learning Dynamic Embeddings from Temporal Interactions." In: **CoRR** abs/1812.02289 (2018).

9. A. Paszke et al. **PyTorch: An Imperative Style, High-Performance Deep Learning Library**. 2019.

10. **Reddit data dump**. http://files.pushshift.io/reddit/.

11. J. W. Pennebaker, M. E. Francis, and R. J. Booth. **Linguistic Inquiry and Word Count**. Mahwah, NJ: Lawerence Erlbaum Associates, 2001.

12. **Wikipedia edit history dump**. https://meta.wikimedia.org/wiki/Data_dumps.

13. B. Hidasi and D. Tikk. "Fast ALS-Based Tensor Factorization for Context-Aware Recommendation from Implicit Feedback." In: **Machine Learning and Knowledge Discovery in Databases**. Ed. by P. A. Flach, T. De Bie, and N. Cristianini. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 67–82.

14. **Kdd cup 2015**. http://moocdata.cn/challenges/kdd-cup-2015.

15. T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár. **Focal Loss for Dense Object Detection**. 2017.